

JAYESH VISHWAKARMA

+91 6264998382 | jayeshvishwakarma6028@gmail.com | [GitHub](#) | [Linkedin](#)

SUMMARY

Final-year MCA student specialising in AI application development and backend engineering. Hands-on experience building RAG pipelines, LLM-integrated REST APIs, and AI agents using FastAPI, LangChain, and GCP. Comfortable owning features end-to-end from prompt design to production deployment. Actively seeking AI/Backend internship or fresher roles.

TECHNICAL SKILLS

Languages:	Python, SQL
AI & LLM:	LangChain, RAG Pipelines, LLM Integration (Gemini, GPT-3.5, LLaMA 3.2), Prompt Engineering, Vector Databases (FAISS), Sentence-Transformers
Backend:	FastAPI, RESTful APIs, Webhook Handling, Async Programming, Docker, Streamlit
Cloud & DevOps:	GCP (Cloud Run, Vertex AI), Git, GitHub
Databases:	PostgreSQL, SQL, RLS (Row-Level Security)
Problem Solving:	150+ DSA problems solved on LeetCode and across platforms

EXPERIENCE

AI/ML Software Developer Intern

Sept 2025 – Dec 2025
CommPlug Innovations, Bangalore

- Built an AI-powered B2B Sales Agent using the Gemini API, integrated with WhatsApp Business API to handle order management conversations end-to-end.
- Implemented Row-Level Security (RLS) in a multi-tenant Django backend, ensuring strict data isolation between B2B clients.
- Designed asynchronous webhook handling with FastAPI and ThreadPoolExecutor, maintaining sub-2-second response times for all incoming messages.
- Wrote automated unit tests and validation pipelines to improve system stability before production deployment.

PROJECTS

PageSense – Context-Aware Browser Assistant [GitHub](#) | [Video Demo](#)

Python · FastAPI · LLaMA 3.2 · FAISS · Sentence-Transformers · Chrome Extension

- Built a Chrome extension chatbot that answers questions about any active webpage using a custom RAG pipeline — extracting live page content, processing it in a FastAPI backend, and returning context-aware answers.
- Integrated FAISS for semantic search with retrieval latency under 2 seconds on pages up to 20k tokens; used Sentence-Transformers for dense embeddings.
- Improved answer relevance on technical documentation through iterative prompt design and chunk-overlap tuning.

InsightLense – Multimodal Research Document Assistant [GitHub](#) | [Video Demo](#)

Python · FastAPI · BM25 · FAISS · Hybrid RAG · Session Persistence

- Architected a multimodal RAG system that parses complex PDF layouts (charts, tables, text) for research document Q&A;
- Built a Hybrid Retrieval pipeline (BM25 + FAISS) combining keyword and semantic search, improving extraction accuracy on structured financial data.
- Implemented fault-tolerant ingestion with rate-limiting and fail-over logic to process 70+ page reports on free-tier API constraints.
- Engineered session persistence and automated cache cleanup, eliminating redundant re-processing for returning users.

AI-Driven Financial Fraud Detection API [GitHub](#) | [Video Demo](#)

Python · XGBoost · FastAPI · GPT-3.5 · Feature Engineering

- Trained a fraud detection model on 6.3M+ financial transactions, optimising for high recall, served via FastAPI with sub-100ms inference.
- Engineered behavioural features capturing anomalies such as accounts drained to zero via unknown transfers.
- Integrated GPT-3.5 to auto-generate model risk reports with explainability metrics, making decisions interpretable for non-technical stakeholders.

EDUCATION

Integrated Master of Computer Applications (MCA)

Acropolis Institute of Technology and Research

Coursework: Data Structures & Algorithms, DBMS, Computer Networks, Operating Systems

Expected 2027
Indore

CURRENTLY EXPLORING

Agentic AI systems • Tool-use patterns with LangGraph • OpenAI Assistants API • Advanced RAG optimisation